

Digital Image Processing

Image Compression - 2

Makarand Tapaswi



Center for Visual Information Technology (CVIT), IIIT Hyderabad



Most slides borrowed from Ravi Kiran @CVIT!

Diversion: Floating point standards

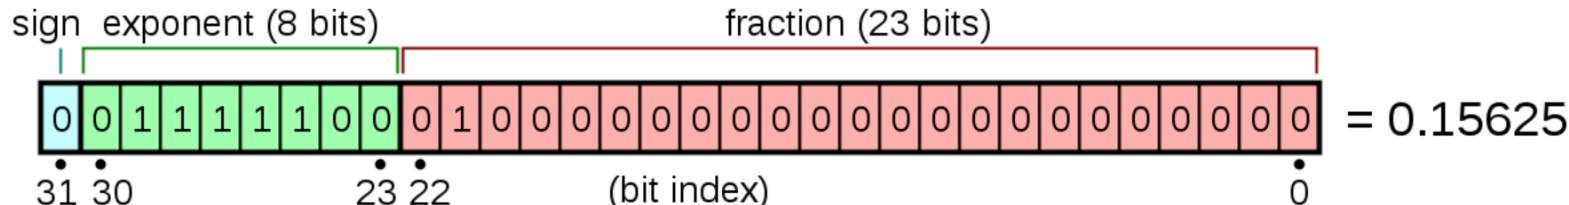
- Pick the number based on suitable exponent

- float32 (IEEE 754)
 - Sign bit: 1 bit
 - Exponent width: 8 bits
 - Significand precision: 23 bits

$$(-1)^{b_{31}} \times 2^{(b_{30}b_{29}\dots b_{23})_2 - 127} \times (1.b_{22}b_{21}\dots b_0)_2,$$

which yields

$$\text{value} = (-1)^{\text{sign}} \times 2^{(E-127)} \times \left(1 + \sum_{i=1}^{23} b_{23-i} 2^{-i} \right).$$



LZW Coding

(addresses interpixel redundancy)

- Requires no prior knowledge of symbol probabilities.
- Assigns **fixed length** code words to **variable length** symbol sequences.
 - There is **no** one-to-one correspondence between source symbols and code words.
- Included in GIF, TIFF and PDF file formats

LZW Coding

- A **codebook** (or **dictionary**) needs to be constructed.
- Initially, the first 256 entries of the dictionary are assigned to the gray levels $0, 1, 2, \dots, 255$ (i.e., assuming 8 bits/pixel)

Consider a 4×4 , 8 bit image

39	39	126	126
39	39	126	126
39	39	126	126
39	39	126	126

Initial Dictionary

Dictionary Location	Entry
0	0
1	1
.	.
255	255
256	-
511	-

LZW Coding (cont'd)

As the encoder examines image pixels, gray level sequences (i.e., blocks) that are not in the dictionary are assigned to a new entry.

39 39 126 126
39 39 126 126
39 39 126 126
39 39 126 126

Dictionary Location	Entry
0	0
1	1
.	.
255	255
256	39-39
511	-

- Is 39 in the dictionary.....Yes
 - What about 39-39.....No
- * Add 39-39 at location 256

Example

39 39 126 126
 39 39 126 126
 39 39 126 126
 39 39 126 126

CR = empty

repeat

P=next pixel

CS=CR + P

If CS is found:

- (1) No Output
- (2) CR=CS

else:

- (1) Output D(CR)
- (2) Add CS to D
- (3) CR=P

Concatenated Sequence: CS = CR + P

(CR) (P)

Currently Recognized Sequence	Pixel Being Processed	Encoded Output	Dictionary Location (Code Word)	Dictionary Entry
	39	39	256	39-39
39	39	39	257	39-126
39	126	126	258	126-126
126	126	126	259	126-39
126	39	256	260	39-39-126
39	39	258	261	126-126-39
39-39	126	260	262	39-39-126-126
126-126	39	259	263	126-39-39
39	126	257	264	39-126-126
126		126		

Decoding LZW

- Use the dictionary for decoding the “encoded output” sequence.
- The dictionary need not be sent with the encoded output.
- Can be built on the “fly” by the decoder as it reads the received code words.

Run-length coding (RLC) (addresses interpixel redundancy)

- Reduce the size of a repeating string of symbols (i.e., runs):

1 1 1 1 1 0 0 0 0 0 0 1 → (1,5) (0, 6) (1, 1)

a a a b b b b b c c → (a,3) (b, 6) (c, 2)

- Encodes a run of symbols into two bytes: **(symbol, count)**
- Can compress any type of data but cannot achieve high compression ratios compared to other compression methods.

Combining Huffman Coding with Run-length Coding

- Assuming that a message has been encoded using Huffman coding, additional compression can be achieved using run-length coding.

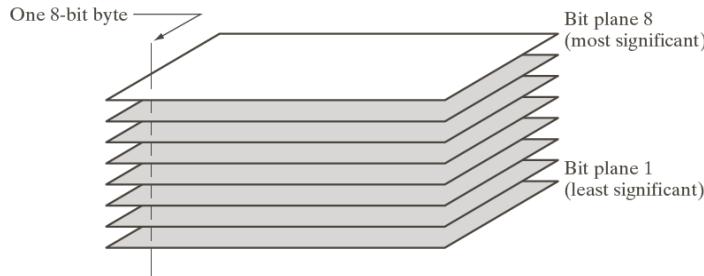
0 1 0 1 0 0 1 1 1 1 0 0

e.g., (0,1)(1,1)(0,1)(1,0)(0,2)(1,4)(0,2)

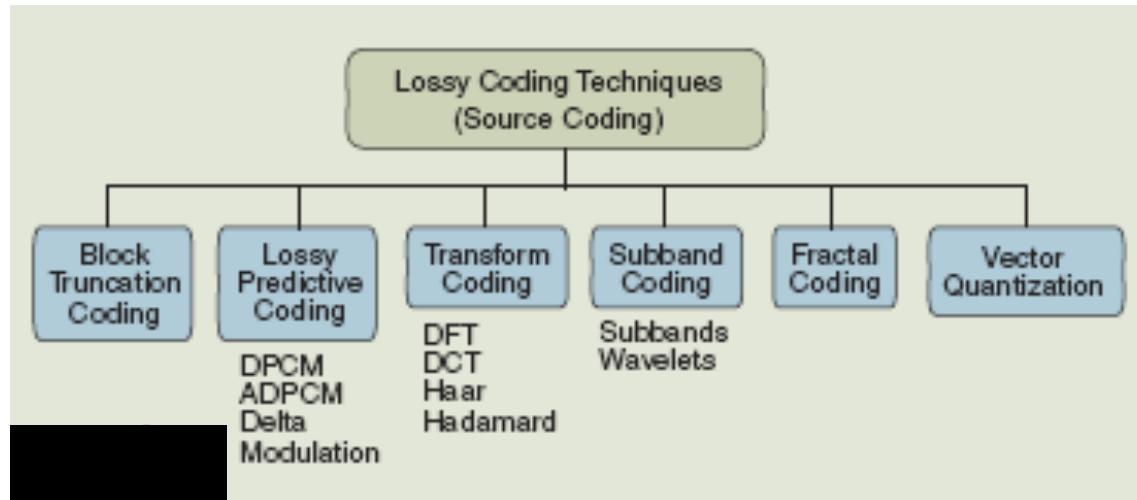
Bit-plane coding

(addresses interpixel redundancy)

- Process each **bit plane** individually.
 - (1) Decompose an image into a series of binary images.
 - (2) Compress each binary image (e.g., using run-length coding)

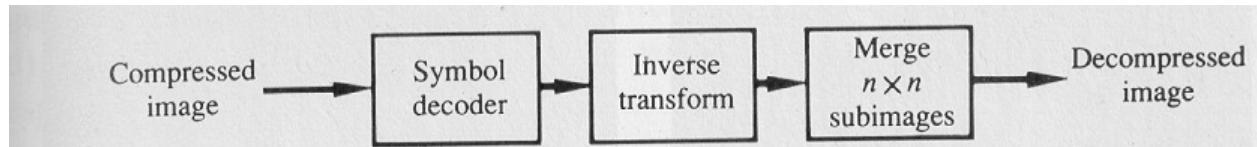
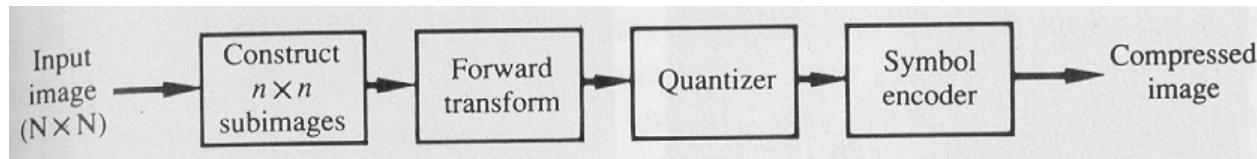


Lossy Methods - Taxonomy



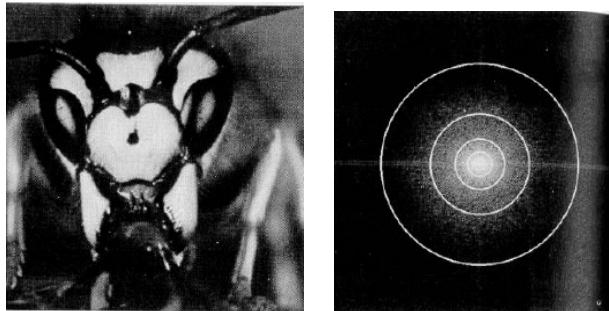
Lossy Compression

- Transform the image into some other domain to reduce interpixel redundancy.



Example: Fourier Transform

$$f(x, y) = \frac{1}{N} \sum_{u=0}^{N-1} \sum_{v=0}^{N-1} F(u, v) e^{\frac{j2\pi(ux+vy)}{N}}, \quad x, y = 0, 1, \dots, N-1$$



Note that the magnitude of the FT decreases, as u, v increase!

K << N

$$\hat{f}(x, y) = \frac{1}{N} \sum_{u=0}^{K-1} \sum_{v=0}^{K-1} F(u, v) e^{\frac{j2\pi(ux+vy)}{N}}, \quad x, y = 0, 1, \dots, N-1$$

$\sum_{x,y} (\hat{f}(x, y) - f(x, y))^2$ is very small !!

Transform Selection

$$f(x, y) = \sum_{u=0}^{N-1} \sum_{v=0}^{N-1} T(u, v) h(x, y, u, v)$$

- $T(u,v)$ can be computed using various transformations, for example:
 - DFT
 - DCT (Discrete Cosine Transform)
 - KLT (Karhunen-Loeve Transformation) or Principal Component Analysis (PCA)
 - DWT (Discrete Wavelet Transform)
- JPEG using DCT for handling interpixel redundancy.

DCT (Discrete Cosine Transform)

[proposed by Nasir Ahmed, T. Natarajan, K.R.Rao (1972)]

Forward: $C(u, v) = \alpha(u)\alpha(v) \sum_{x=0}^{N-1} \sum_{y=0}^{N-1} f(x, y) \cos\left(\frac{(2x+1)u\pi}{2N}\right) \cos\left(\frac{(2y+1)v\pi}{2N}\right),$

$u, v=0,1,\dots,N-1$

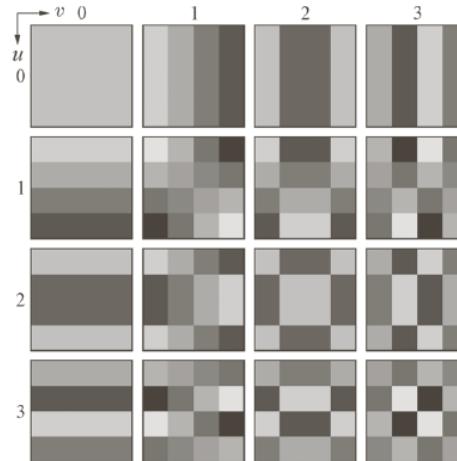
Inverse: $f(x, y) = \sum_{u=0}^{N-1} \sum_{v=0}^{N-1} \alpha(u)\alpha(v) C(u, v) \cos\left(\frac{(2x+1)u\pi}{2N}\right) \cos\left(\frac{(2y+1)v\pi}{2N}\right),$

$x, y=0,1,\dots,N-1$

$$\alpha(u) = \begin{cases} \sqrt{1/N} & \text{if } u=0 \\ \sqrt{2/N} & \text{if } u>0 \end{cases}$$
$$\alpha(v) = \begin{cases} \sqrt{1/N} & \text{if } v=0 \\ \sqrt{2/N} & \text{if } v>0 \end{cases}$$

DCT (cont'd)

- Basis functions for a 4x4 image (i.e., cosines of different frequencies).



DCT (cont'd)

Using
8 x 8 sub-images
yields 64 coefficients
per sub-image.

Reconstructed images
by truncating
50% of the
coefficients

More compact
transformation

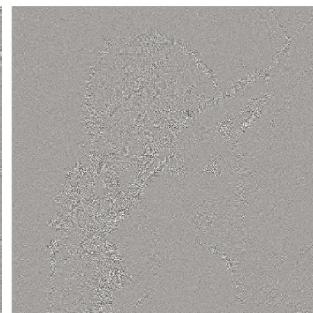
DFT



WHT



DCT



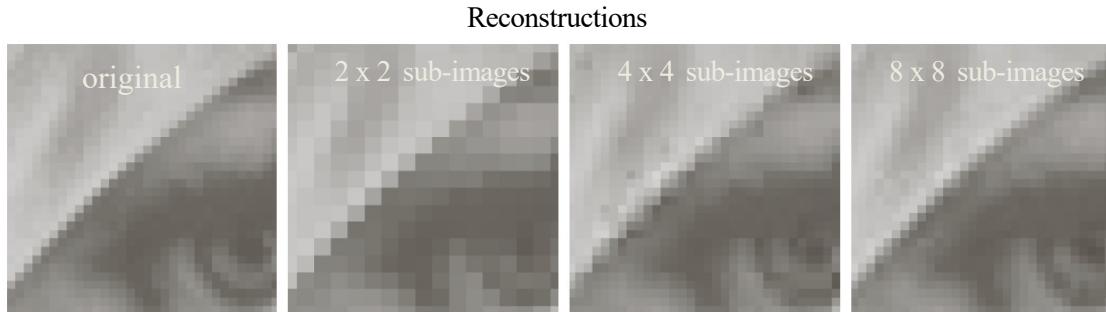
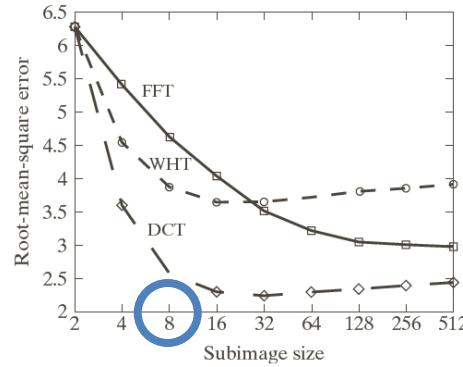
RMS error: 2.32

1.78

1.13

DCT (cont'd)

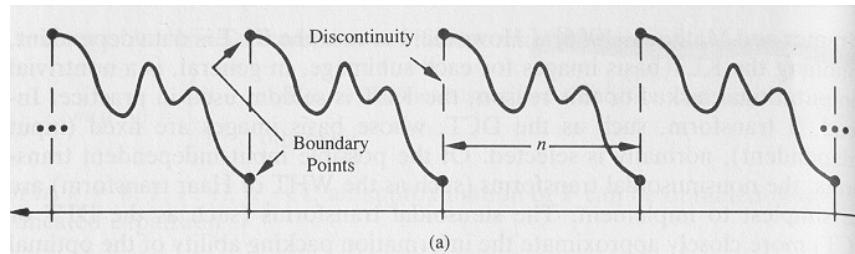
- Sub-image size selection



DCT (cont'd)

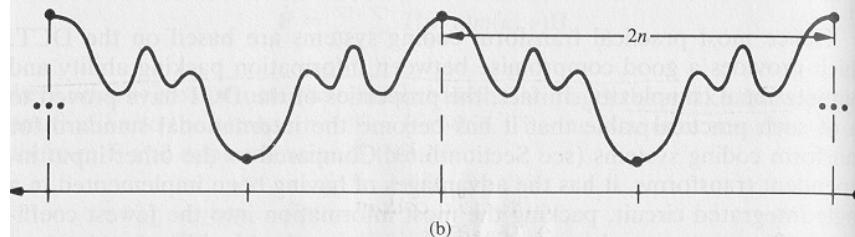
- DCT minimizes "blocking artifacts" (i.e., boundaries between subimages do not become very visible).

DFT
has **n-point** periodicity



(a)

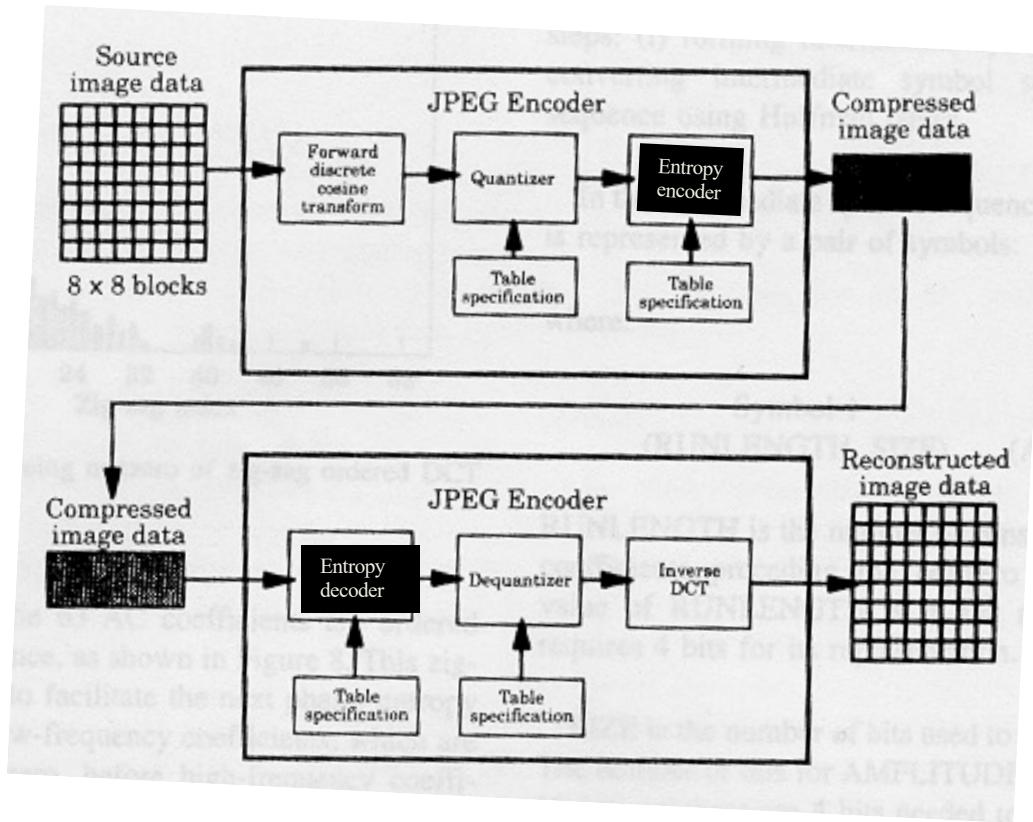
DCT
has **2n-point** periodicity



(b)

JPEG Compression

Accepted as an international image compression standard in 1992.



JPEG - Steps

1. Divide image into 8x8 subimages.

For each subimage do:

2. Shift the gray-levels in the range [-128, 127]

3. Apply DCT → 64 coefficients

1 DC coefficient: $F(0,0)$

63 AC coefficients: $F(u,v)$

Example

(a) Original 8x8 block	(b) Shifted block	(c) Block after FDCT Eqn. (5)
140 144 147 1140 140 155 179 175 144 152 140 147 140 148 167 179 152 155 136 167 163 182 152 172 168 145 156 160 152 155 136 160 162 148 156 148 140 136 147 162 147 167 140 155 155 140 136 162 136 156 123 167 162 144 140 147 148 155 136 155 152 147 147 136	12 16 19 12 11 27 51 47 16 24 12 19 12 20 39 51 24 27 8 39 35 34 24 44 40 17 28 32 24 27 8 32 34 20 28 20 12 8 19 34 19 39 12 27 27 12 8 34 8 28 -5 39 34 16 12 19 20 27 8 27 24 19 19 8	185 -17 14 -8 23 -9 -13 -18 20 -34 26 -9 -10 10 13 6 -10 -23 -1 6 -18 3 -20 0 -8 -5 14 -14 -8 -2 -3 8 -3 9 7 1 -11 17 18 15 3 -2 -18 8 8 -3 0 -6 8 0 -2 3 -1 -7 -1 -1 0 -7 -2 1 1 4 -6 0

[-128, 127]

(non-centered
spectrum)

JPEG Steps

4. Quantize the coefficients (i.e., reduce the amplitude of coefficients that do not contribute a lot).

$$C_q(u, v) = \text{Round}\left[\frac{C(u, v)}{Q(u, v)}\right]$$


$Q(u, v)$: quantization table

Example

- Quantization Table $Q[i][j]$

```
for i=0 to n;  
  for j=0 to n;  
    Q[i,j] = 1 + (1+i+j)*quality;  
  end j;  
end i;
```

(d) Quantization table
(quality = 2)

3	5	7	9	11	13	15	17
5	7	9	11	13	15	17	19
7	9	11	13	15	17	19	21
9	11	13	15	17	19	21	23
11	13	15	17	19	21	23	25
13	15	17	19	21	23	25	27
15	17	19	21	23	25	27	29
17	19	21	23	25	27	29	31

$$1 \leq quality \leq 25$$

(best - low compression)

(worst - high compression)

Example (cont'd)

(c) Block after FDCT
Eqn. (5)

185	-17	14	-8	23	-9	-13	-18
20	-34	26	-9	-10	10	13	6
-10	-23	-1	6	-18	3	-20	0
-8	-5	14	-14	-8	-2	-3	8
-3	9	7	1	-11	17	18	15
3	-2	-18	8	8	-3	0	-6
8	0	-2	3	-1	-7	-1	-1
0	-7	-2	1	1	4	-6	0

(d) Quantization table
(quality = 2)

3	5	7	9	11	13	15	17
5	7	9	11	13	15	17	19
7	9	11	13	15	17	19	21
9	11	13	15	17	19	21	23
11	13	15	17	19	21	23	25
13	15	17	19	21	23	25	27
15	17	19	21	23	25	27	29
17	19	21	23	25	27	29	31

Quantization



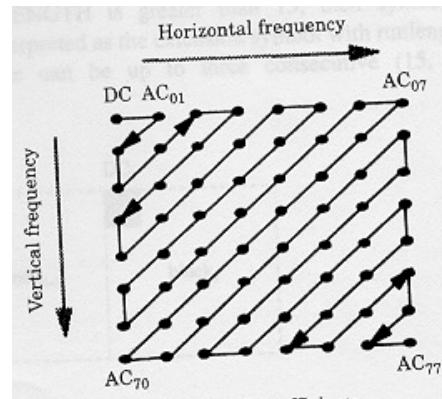
(e) Block after quantization
Eqn. (6)

JPEG Steps (cont'd)

5. Order the coefficients using **zig-zag** ordering

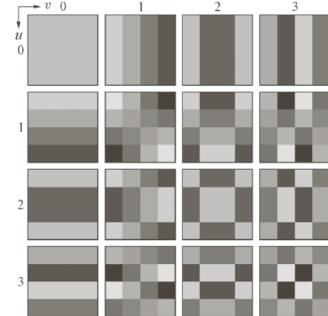
- Creates long runs of zeros (i.e., ideal for run-length encoding)

(e) Block after quantization
Eqn. (6)



(f) Zig-zag sequence

61,-3,4,-1,-4,2,0,2,-2,0,0,0,0,0,2,0,0,0,1,0,0,0,0,0,0,-1,0,0,-1,0,0,
0,0,-1,0,0,0,0,0,0,-1,0



JPEG Steps (cont'd)

6. Encode coefficients:

- 6.1 Form “**intermediate**” symbol sequence.
- 6.2 Encode “**intermediate**” symbol sequence into a binary sequence.

Intermediate Symbol Sequence – DC coeff

(f) Zig-zag sequence

1

(g) Intermediate symbol sequence

$$(6)(61), (0,2)(-3), (0,3)(4), (0,1)(-1), (0,3)(-4), (0,2)(2), (1,2)(2), (0,2)(-2), (5,2)(2), (3,1)(1), (6,1)(-1), (2,1)(-1), (4,1)(-1), (7,1)(-1), (0,0)$$

symbol_1 (SIZE)

symbol_2 (AMPLITUDE)

DC

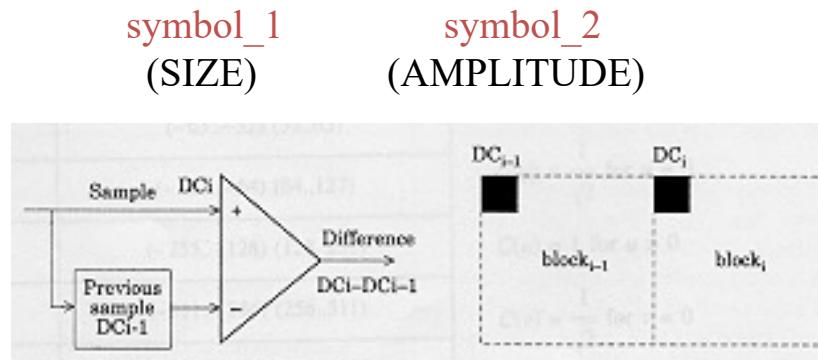
(6)

(61)

SIZE: # bits need to encode the coefficient

DC Coefficient Encoding

**predictive
coding:**



Intermediate Symbol Sequence – AC coeff

(f) Zig-zag sequence

61,-3,4,-1,-4,2,0,2,-2,0,0,0,0,0,2,0,0,0,1,0,0,0,0,0,0,-1,0,0,-1,0,0,
0,0,-1,0,0,0,0,0,0,0,-1,0

(g) Intermediate symbol sequence

$$(6)(61), (0,2)(-3), (0,3)(4), (0,1)(-1), (0,3)(-4), (0,2)(2), (1,2)(2), (0,2)(-2) \\ \text{---} (5,2)(2), (3,1)(1), (6,1)(-1), (2,1)(-1), (4,1)(-1), (7,1)(-1), (0,0)$$

symbol 1 (RUN-LENGTH, SIZE)

symbol 2 (AMPLITUDE)

end of block

AC

(0, 2)

(-3)

RUN-LENGTH: run of zeros preceding coefficient
SIZE: # bits for encoding the amplitude of coefficient

Note: If RUN-LENGTH > 15, use symbol (15,0) .

Example: AC Coefficients Encoding

Symbol_1

(Variable Length Code (VLC))

(Runlength, size)	Code word
(0,,0) EOB	1010
(0,1)	00
(0,2)	01
(0,3)	100
(1,2)	11011
(2,1)	11100
(3,1)	111010
(4,1)	111011
(5,2)	1111110111
(6,1)	11111011
(7,1)	11111010

Symbol_2

(Variable Length Integer (VLI))

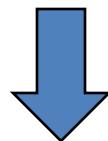
# bits	Amplitude range
1	(-1, 1)
2	(-3, -2) (2,3)
3	(-7, -4) (4..7)
4	(-15..-8) (8..15)
5	(-31..-16) (16..31)
6	(-63..-32) (32..63)
7	(-127..-64) (64..127)
8	(-255..-128) (128..255)
9	(-511..-256) (256..511)
10	(-1023..-512) (512..1023)

(1,4) (12) → (111110110 1100)
VLC VLI

Final Symbol Sequence

(g) Intermediate symbol sequence

(6)(61),(0,2)(-3),(0,3)(4),(0,1)(-1),(0,3)(-4),(0,2)(2),(1,2)(2),(0,2)(-2),
[REDACTED] (5,2)(2),(3,1)(1),(6,1)(-1),(2,1)(-1),(4,1)(-1),(7,1)(-1),(0,0)



(e) Encoded bit sequence (total 98 bits)

111011110100100100000100011011011011100101111111011
11011101011111011011100011101101111101001010

What is the effect of the “Quality” parameter?



(58k bytes)



(21k bytes)



(8k bytes)

lower compression

higher compression

$$1 \leq \text{quality} \leq 25$$



ONE DOES NOT SIMPLY

USE HIGH LEVEL OF JPEG COMPRESSION

Case Study: Fingerprint Compression

- FBI is digitizing fingerprints at 500 dots per inch with 8 bits of grayscale resolution.
- A single fingerprint card turns into about 10 MB of data!



A sample fingerprint image
 768×768 pixels = 589,824 bytes

WSQ Fingerprint Compression

- An image coding standard for digitized fingerprints employing the **Discrete Wavelet Transform (Wavelet/Scalar Quantization or WSQ)**.
- Developed and maintained by:
 - FBI
 - Los Alamos National Lab (LANL)
 - National Institute for Standards and Technology (NIST)

Need to Preserve Fingerprint Details



The "white" spots in the middle of the black ridges are *sweat pores* and they are admissible points of identification in court.

These details are just a couple pixels wide!

What compression scheme should be used?

- Lossless or lossy compression?
- In practice lossless compression methods haven't done better than **2:1** on fingerprints!
- Does JPEG work well for fingerprint compression?

Results using JPEG compression

file size 45853 bytes

compression ratio: 12.9



Fine details have been lost.

Image has an artificial “**blocky**” pattern superimposed on it.

Artifacts will affect the performance of fingerprint recognition.

Results using WSQ compression

file size 45621 bytes

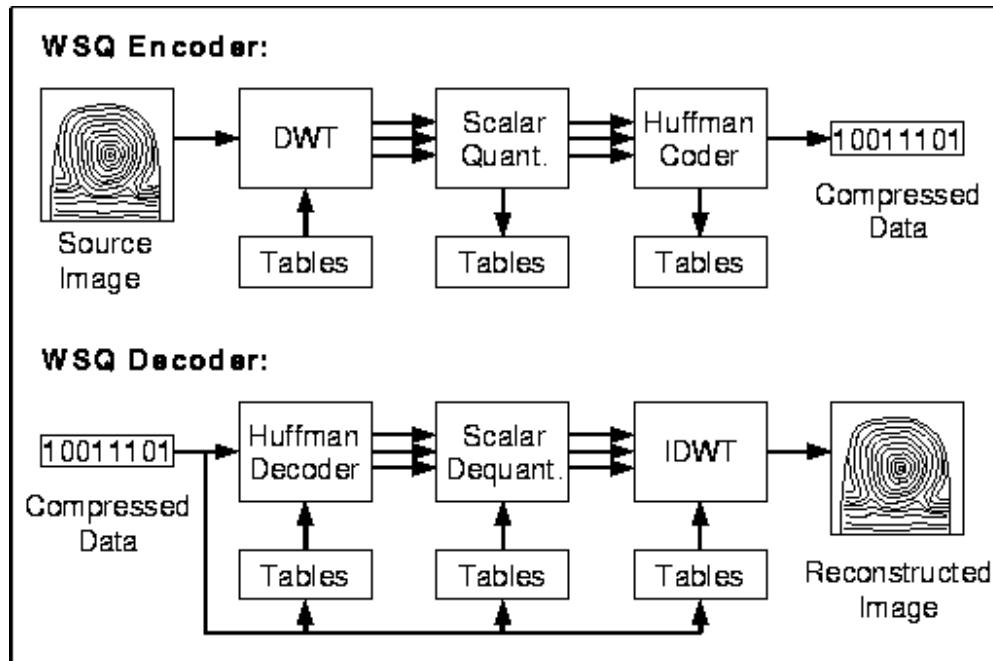
compression ratio: 12.9



Fine details are better
preserved.

No “blocky” artifacts.

WSQ Algorithm



Compression ratio

- FBI's target bit rate is around 0.75 bits per pixel (bpp)
- This corresponds to a compression ratio of $8/0.75=10.7$
- Target bit rate can set via a parameter, similar to the "quality" parameter in JPEG.

Varying compression ratio (cont'd)

Original image 768 x 768 pixels (589824 bytes)



Varying compression ratio (cont'd)

0.9 bpp compression

WSQ image, file size 47619 bytes,
compression ratio 12.4



JPEG image, file size 49658 bytes,
compression ratio 11.9



Varying compression ratio (cont'd)

0.75 bpp compression

WSQ image, file size 39270 bytes,
compression ratio 15.0



JPEG image, file size 40780 bytes,
compression ratio 14.5



Varying compression ratio (cont'd)

0.6 bpp compression

WSQ image, file size 30987 bytes,
compression ratio 19.0



JPEG image, file size 30081 bytes,
compression ratio 19.6

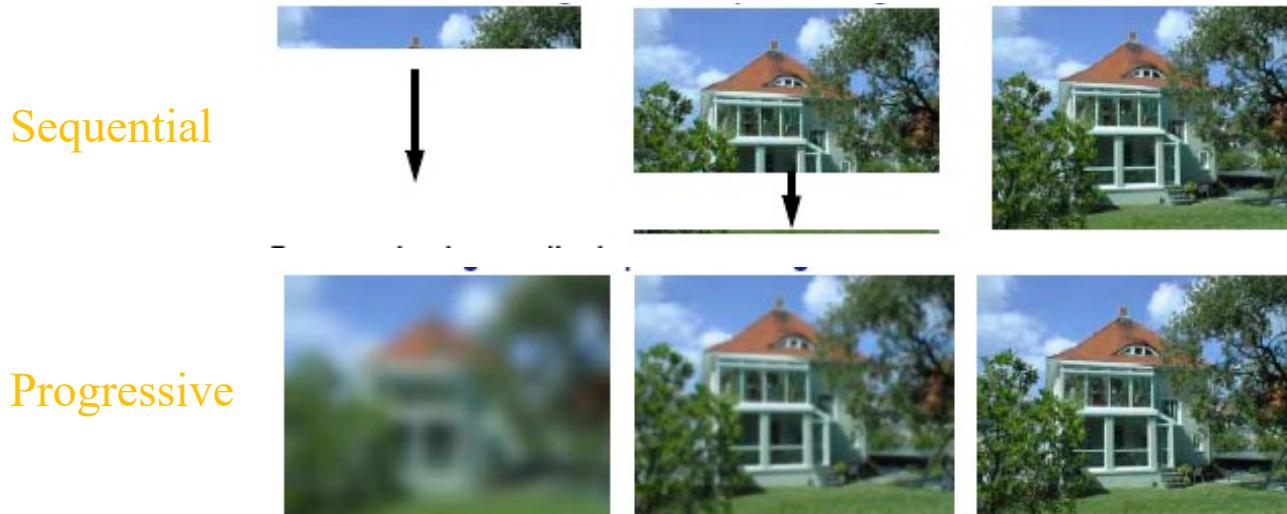


JPEG Modes

- JPEG supports several different modes
 - Sequential Mode
 - Progressive Mode
 - Hierarchical Mode
 - Lossless Mode
- The default mode is “sequential”
 - Image is encoded in a **single scan** (left-to-right, top-to-bottom).

Progressive JPEG

- Image is encoded in **multiple scans**, in order to produce a quick, rough decoded image when transmission time is long.



Color Images



Y

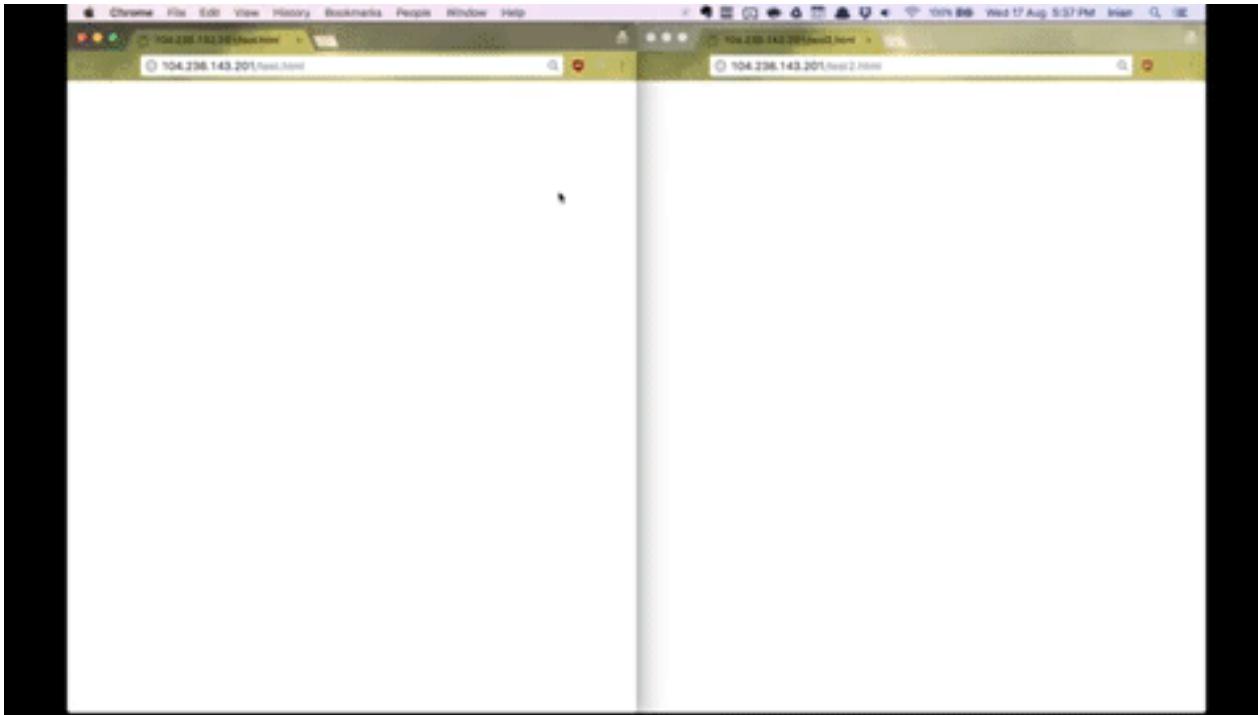


Cb



Cr

- Different quantization matrices for chrominance and luminance
- Chroma subsampling (use reduced resolution of chroma channels)



JPEG at 0.125 bpp (enlarged)



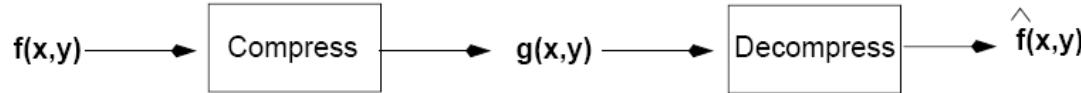
JPEG2000 at 0.125 bpp



Other popular formats

- GIF → lossy
- PNG → lossless
- Video
 - MPEG etc. (exploit temporal redundancy also)

Fidelity Criteria



- How close is $f(x,y)$ to $\hat{f}(x,y)$?
- Criteria
 - Subjective: based on human observers
 - Objective: mathematically defined criteria

Subjective Fidelity Criteria

Value	Rating	Description
1	Excellent	An image of extremely high quality, as good as you could desire.
2	Fine	An image of high quality, providing enjoyable viewing. Interference is not objectionable.
3	Passable	An image of acceptable quality. Interference is not objectionable.
4	Marginal	An image of poor quality; you wish you could improve it. Interference is somewhat objectionable.
5	Inferior	A very poor image, but you could watch it. Objectionable interference is definitely present.
6	Unusable	An image so bad that you could not watch it.

Quality measurement: judged by human viewers

- Five scale system on the degree of impairment
 1. Impairment is not noticeable
 2. Impairment is just noticeable
 3. Impairment is definitely noticeable, but not objectionable
 4. Impairment is objectionable
 5. Impairment is extremely objectionable

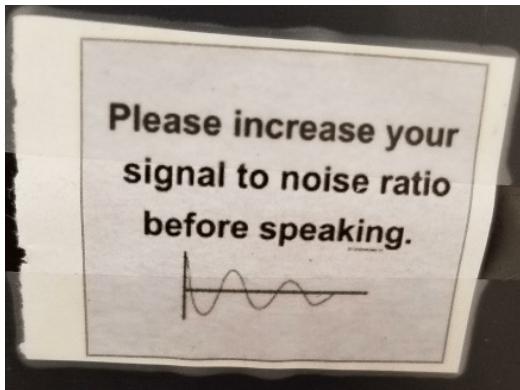
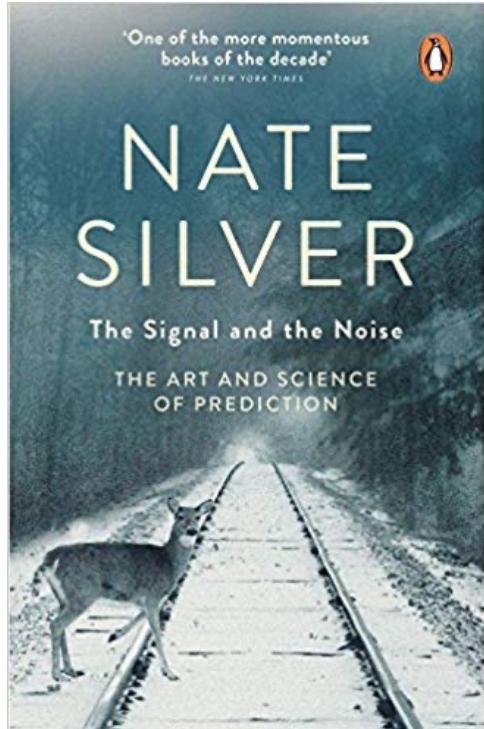
Advantages: relies on HVS

Drawbacks: time, viewing conditions, viewers?



Quality measurement: Signal to noise ratio

$e(x, y)$



1

$$SNR_{ms} = 10 \log_{10} \left(\frac{\sum_{x=0}^{M-1} \sum_{y=0}^{N-1} g(x, y)^2}{MN \cdot E_{ms}} \right)$$

$$E_{ms} = \frac{1}{MN} \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} e(x, y)^2$$

Decibel
(dB)

Quality measurement: Signal to noise ratio

$$e(x, y) = f(x, y) - g(x, y). \quad E_{\text{ms}} = \frac{1}{MN} \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} e(x, y)^2$$

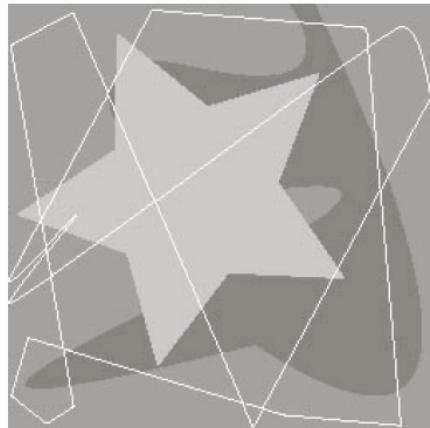
1 $SNR_{ms} = 10 \log_{10} \left(\frac{\sum_{x=0}^{M-1} \sum_{y=0}^{N-1} g(x, y)^2}{MN \cdot E_{\text{ms}}} \right)$

2 $PSNR = 10 \log_{10} \left(\frac{255^2}{E_{\text{ms}}} \right)$

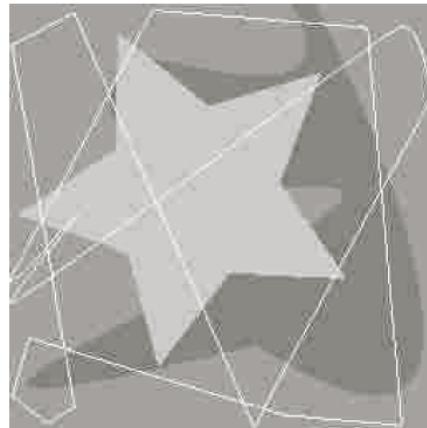
**Decibel
(dB)**

Subjective vs Objective Fidelity Criteria

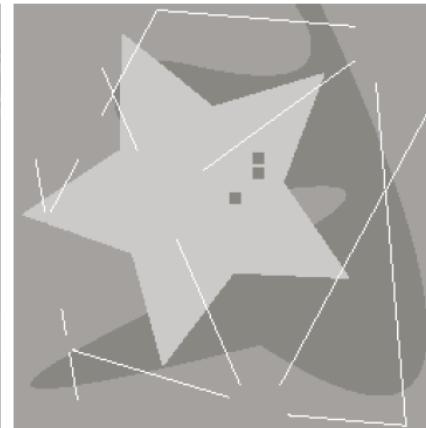
RMSE = 5.17



RMSE = 15.67



RMSE = 14.17



Next generation tech: ‘Semantic’ data compression !



The screenshot shows a Facebook photo interface. At the top right are 'Tag Photo' and 'Edit' buttons. Below them are 'Like', 'Comment', and 'Share' buttons, followed by a list of users who liked the post ('Suresh Mittal and 7 others'). A text input field says 'Write a comment...'. At the bottom is a 'Suggested Pages' section.

Elements Console Sources Network Timeline Profiles Application Security Audits

```
<div class="stage" data-ft="{"tn":"E"}>
  <div class="fbPhotosPhotoTagboxes tagContainer" id="fbPhotoSnowliftTagBoxes"></div>
  <div class="fbPhotoTagApproval hidden_elem" id="fbPhotoSnowliftTagApproval"></div>
  <div class="5ba1" id="fbPhotoSnowliftComputerVisionInfo"></div>
<div class="2-sx" style="width: 100%; height: 455px; >
  
<div class="videoStage" data-ft="{"tn":"F"}></div>
```

Styles Co
Filter
element.s
}
.4d3w .s
img.spotl
vertical
}



The screenshot shows a Facebook post from 'Stephanie Hughes' posted 'Yesterday at 10:14 PM'. The post features a large pepperoni and olive pizza. A callout bubble in the bottom right corner says 'Image may contain: pizza, food'. The post has '48 Likes' and '6 Comments'. At the top right are 'Like', 'Comment', and 'Share' buttons.

Humans are still the best lossy image compressors

Ashutosh Bhowm^{1,*}, Soham Mukherjee^{2,*}, Sean Yang^{3,*},
Shubham Chandak⁴, Irena Fischer-Hwang⁴, Kedar Tatwawadi⁴, Tsachy Weissman⁴

¹Palo Alto High School

²Monta Vista High School

³Saint Francis High School

⁴Stanford University

schandak@stanford.edu

Abstract

Lossy image compression has been studied extensively in the context of typical loss functions such as RMSE, MS-SSIM, etc. However, it is not well understood what loss function might be most appropriate for human perception. Furthermore, the availability of massive public image datasets appears to have hardly been exploited in image compression. In this work, we perform compression experiments in which one human describes images to another, using publicly available images and text instructions. These image reconstructions are rated by human scorers on the Amazon Mechanical Turk platform and compared to reconstructions obtained by existing image compressors. In our experiments, the humans outperform the state of the art compressor WebP in the MTurk survey on most images, which shows that there is significant room for improvement in image compression for human perception.

Data: The images, results and additional data is available at <https://compression.stanford.edu/human-compression>.

Introduction

Since the advent of electronic media, image compression has been studied extensively, leading to multiple image formats and compression techniques such as PNG [1], JPEG [2], JPEG2000 [3], JPEG XR [4], BPG [5] and WebP [6]. In order to

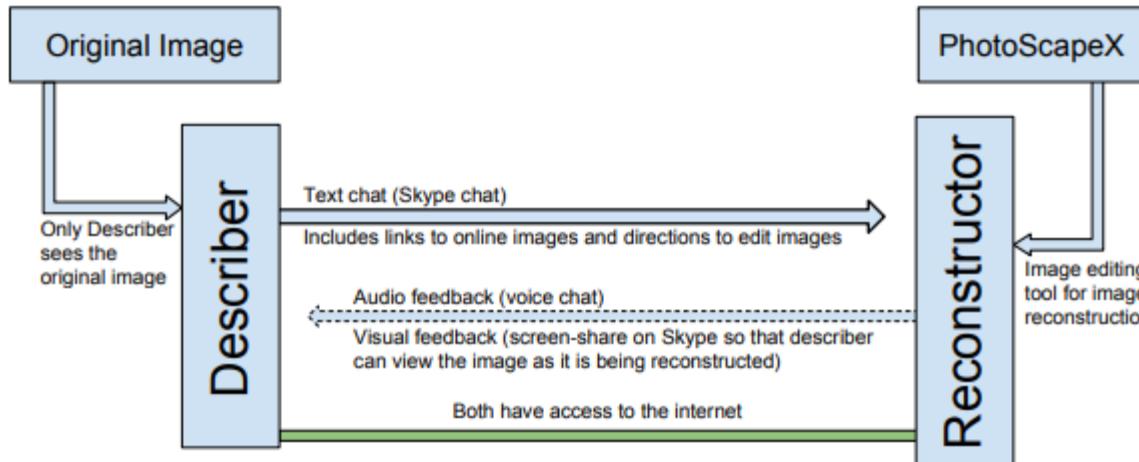
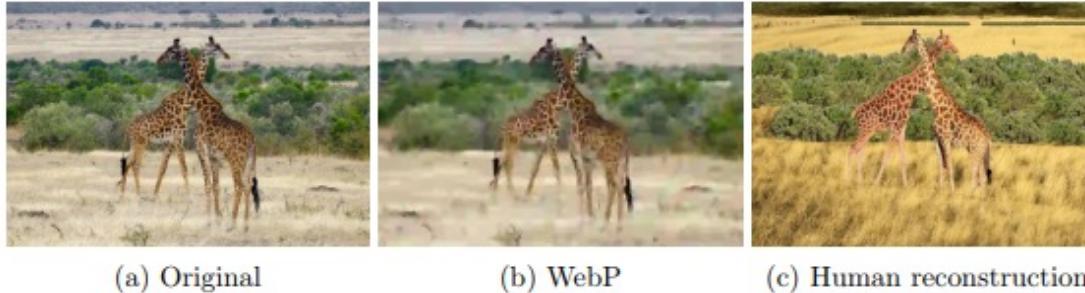


Figure 2: Block diagram showing the human compression process. The describer attempts to describe the image with URL links and text instructions. The describer can see the image as it is being reconstructed and can also hear the reconstructor's voice.



<https://www.worldwildlife.org/habitats/grasslands>



try transformations
elongate the fence bit
only focus on the vertical...



there's a line of shrubbery that goes across the middle third of the image...
that's the largest bush in the pic
so keep the others sizes equal to or smaller than that and make it look
continuous
and make sure to make the bushes smaller as you work your way up so
that there's a sense of depth...



Try and make the grass look less tall on the bottom...



when you're done with that take a look at these
https://public-media.smithsonianmag.com/filer/32/f2/32f24473-b380-43f5-94df-da0e58644439/16301090250_acf80be87f_o.jpg
<https://img.purch.com/w/192/aHR0cDovL3d3dy5saXZlc2NpZW5jZS5jb20vaW1hZ2VzL2kvMDAwLzA2OC8wOTQvaTMwMC9naXJhZmZlLmpwZz8xNDA1MDA4NDQy>
sure
while you're editing that giraffe
its spots are too dark
make it look like the other giraffe...

make the right one bigger than the left
make the heads level
wait back
put the left one where it was before
good
now move the right giraffe to the left so that their necks cross
good
move them both to the center
make them both taller as well
their heads should be above the middle line of shrubs...



Final reconstruction



Original image



📅 15. JUNE 📍 SEATTLE, WASHINGTON

NTIRE 2020

New Trends in Image Restoration and Enhancement workshop
and challenges on image and video restoration and enhancement

in conjunction with [CVPR 2020](#)

Reference

- Ch 8, G&W textbook