


Modelling Series

- In many situations one must consider a *series* of inputs to produce an output
 - Outputs too may be a series
- Examples: ..

What did I say?

“To be” or not “to be”??



- Speech Recognition
 - Analyze a series of spectral vectors, determine what was said
- Note: Inputs are sequences of vectors. Output is a classification result

What is he talking about?

“Football” or “basketball”?




The Steelers, meanwhile, continue to struggle to make stops on defense. They've allowed, on average, 30 points a game, and have shown no signs of improving anytime soon.

- Text analysis
 - E.g. analyze document, identify topic
 - Input series of words, output classification output
 - E.g. read English, output French
 - Input series of words, output series of words

Should I invest..

To invest or not to invest?




- Note: Inputs are sequences of vectors. Output may be scalar or vector
 - Should I invest, vs. should I not invest in X?
 - Decision must be taken considering how things have fared over time

These are classification and prediction problems


- Consider a sequence of inputs
 - Input vectors
- Produce one or more outputs
- This can be done with neural networks
 - Obviously

Representational shortcut




- Input at each time is a *vector*
- Each layer has many neurons
 - Output layer too may have many neurons
- But will represent everything by simple boxes
 - Each box actually represents an entire *layer with many units*

Representational shortcut



- Input at each time is a *vector*
- Each layer has many neurons
 - Output layer too may have many neurons
- But will represent everything by simple boxes
 - Each box actually represents an entire *layer with many units*


Representational shortcut



- Input at each time is a *vector*
- Each layer has many neurons
 - Output layer too may have many neurons
- But will represent everything as simple boxes
 - Each box actually represents an entire *layer with many units*


The stock prediction problem...

To invest or not to invest?




- Stock market
 - Must consider the series of stock values in the past several days to decide if it is wise to invest today
 - Ideally consider *all* of history

The stock predictor network




- The sliding predictor
 - Look at the last few days
 - This is just a convolutional neural net applied to series data
 - Also called a *Time-Delay neural network*

The stock predictor network




- The sliding predictor
 - Look at the last few days
 - This is just a convolutional neural net applied to series data
 - Also called a *Time-Delay neural network*

The stock predictor network




- The sliding predictor
 - Look at the last few days
 - This is just a convolutional neural net applied to series data
 - Also called a *Time-Delay neural network*

The stock predictor network



- The sliding predictor
 - Look at the last few days
 - This is just a convolutional neural net applied to series data
 - Also called a *Time-Delay neural network*

The stock predictor network




- The sliding predictor
 - Look at the last few days
 - This is just a convolutional neural net applied to series data
 - Also called a *Time-Delay neural network*

Finite-response model

- This is a *finite response* system
 - Something that happens *today* only affects the output of the system for N days into the future
 - N is the *width* of the system

$$Y_t = f(X_t, X_{t-1}, \dots, X_{t-N})$$


The stock predictor



- This is a *finite response* system
 - Something that happens *today* only affects the output of the system for N days into the future
 - N is the *width* of the system

$$Y_t = f(X_t, X_{t-1}, \dots, X_{t-N})$$


The stock predictor



- This is a *finite response* system
 - Something that happens *today* only affects the output of the system for N days into the future
 - N is the *width* of the system

$$Y_t = f(X_t, X_{t-1}, \dots, X_{t-N})$$


The stock predictor



- This is a *finite response* system
 - Something that happens *today* only affects the output of the system for N days into the future
 - N is the *width* of the system

$$Y_t = f(X_t, X_{t-1}, \dots, X_{t-N})$$


The stock predictor



- This is a *finite response* system
 - Something that happens *today* only affects the output of the system for N days into the future
 - N is the *width* of the system

$$Y_t = f(X_t, X_{t-1}, \dots, X_{t-N})$$


The stock predictor



- This is a *finite response* system
 - Something that happens *today* only affects the output of the system for N days into the future
 - N is the *width* of the system

$$Y_t = f(X_t, X_{t-1}, \dots, X_{t-N})$$


The stock predictor



- This is a *finite response* system
 - Something that happens *today* only affects the output of the system for N days into the future
 - N is the *width* of the system


$$Y_t = f(X_t, X_{t-1}, \dots, X_{t-N})$$

Finite-response model



- Something that happens *today* only affects the output of the system for N days into the future
 - **Predictions consider N days of history**
- To consider more of the past to make predictions, you must increase the “history” considered by the system

Finite-response




- Problem: Increasing the “history” makes the network more complex
 - No worries, we have the CPU and memory
 - Or do we?

Systems often have long-term dependencies



- Longer-term trends –
 - Weekly trends in the market
 - Monthly trends in the market
 - Annual trends
 - Though longer historic trends to affect us less than more recent events..

We want *infinite* memory



- Required: *Infinite* response systems
 - What happens today can continue to affect the output forever
 - Possibly with weaker and weaker influence


$$Y_t = f(X_t, X_{t-1}, \dots, X_{t-\infty})$$

Examples of infinite response systems

$$Y_t = f(X_t, Y_{t-1})$$


- Required: Define initial state: Y_{-1} for $t = 0$
 - An input at X_0 at $t = 0$ produces Y_0
 - Y_0 produces Y_1 which produces Y_2 and so on until Y_∞ *even if* $X_1 \dots X_\infty$ are 0
 - i.e. even if there are no further inputs!
 - **A single input influences the output for the rest of time!**
- This is an instance of a NARX network
 - “nonlinear autoregressive network with exogenous inputs”
 - $Y_t = f(X_{0:t}, Y_{0:t-1})$
 - *Output* contains information about the entire past

A one-tap NARX network




- A NARX net with recursion from the output

A one-tap NARX network




- A NARX net with recursion from the output

A one-tap NARX network




- A NARX net with recursion from the output

A one-tap NARX network




- A NARX net with recursion from the output

A one-tap NARX network




- A NARX net with recursion from the output

A one-tap NARX network




- A NARX net with recursion from the output

A one-tap NARX network




- A NARX net with recursion from the output

A one-tap NARX network



- A NARX net with recursion from the output

A more complete representation



- A NARX net with recursion from the output
- Showing all computations
- All columns are identical
- *An input at $t=0$ affects outputs forever*

NARX Networks

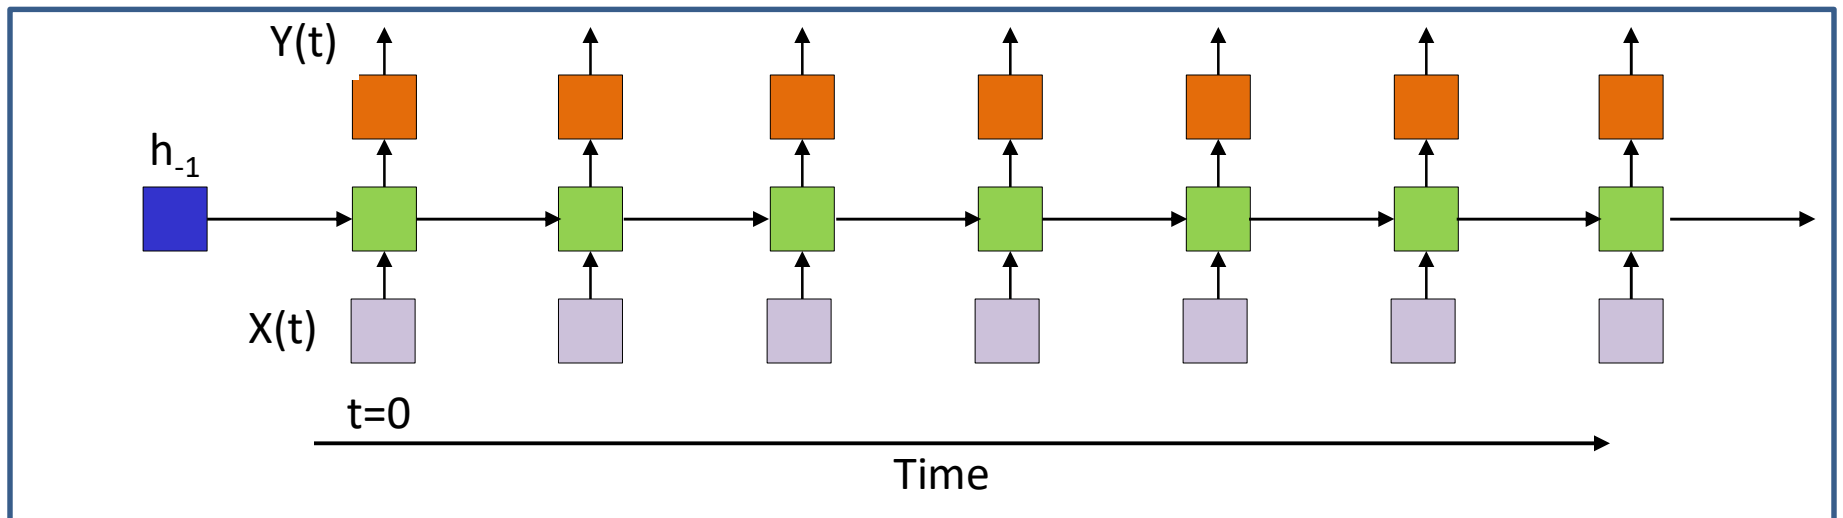

- Very popular for time-series prediction
 - Weather
 - Stock markets
 - As alternate system models in tracking systems
- Any phenomena with distinct “innovations” that “drive” an output
- Note: here the “memory” of the past is in the output itself, and not in the network

An alternate model for infinite response systems: **the state-space model**

$$\begin{aligned}h_t &= f(x_t, h_{t-1}) \\ y_t &= g(h_t)\end{aligned}$$

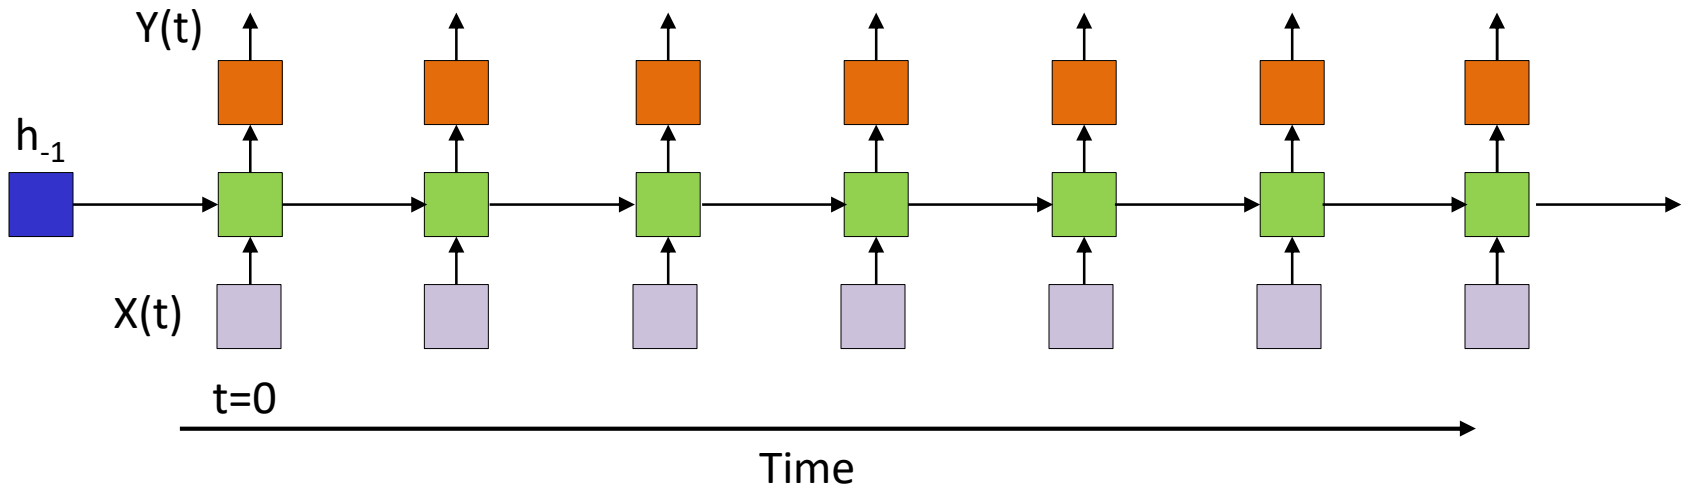
- h_t is the *state* of the network
 - *State* summarizes information about the entire past
 - Model directly embeds the memory in the state
- Need to define initial state h_{-1}
- This is a *fully recurrent* neural network
 - Or simply a *recurrent neural network*

The simple state-space model




- The state (green) at any time is determined by the input at that time, and the state at the previous time
- *An input at $t=0$ affects outputs forever*
- Also known as a recurrent neural net

Single hidden layer RNN




- Recurrent neural network
- All columns are identical
- *An input at $t=0$ affects outputs forever*

Multiple recurrent layer RNN



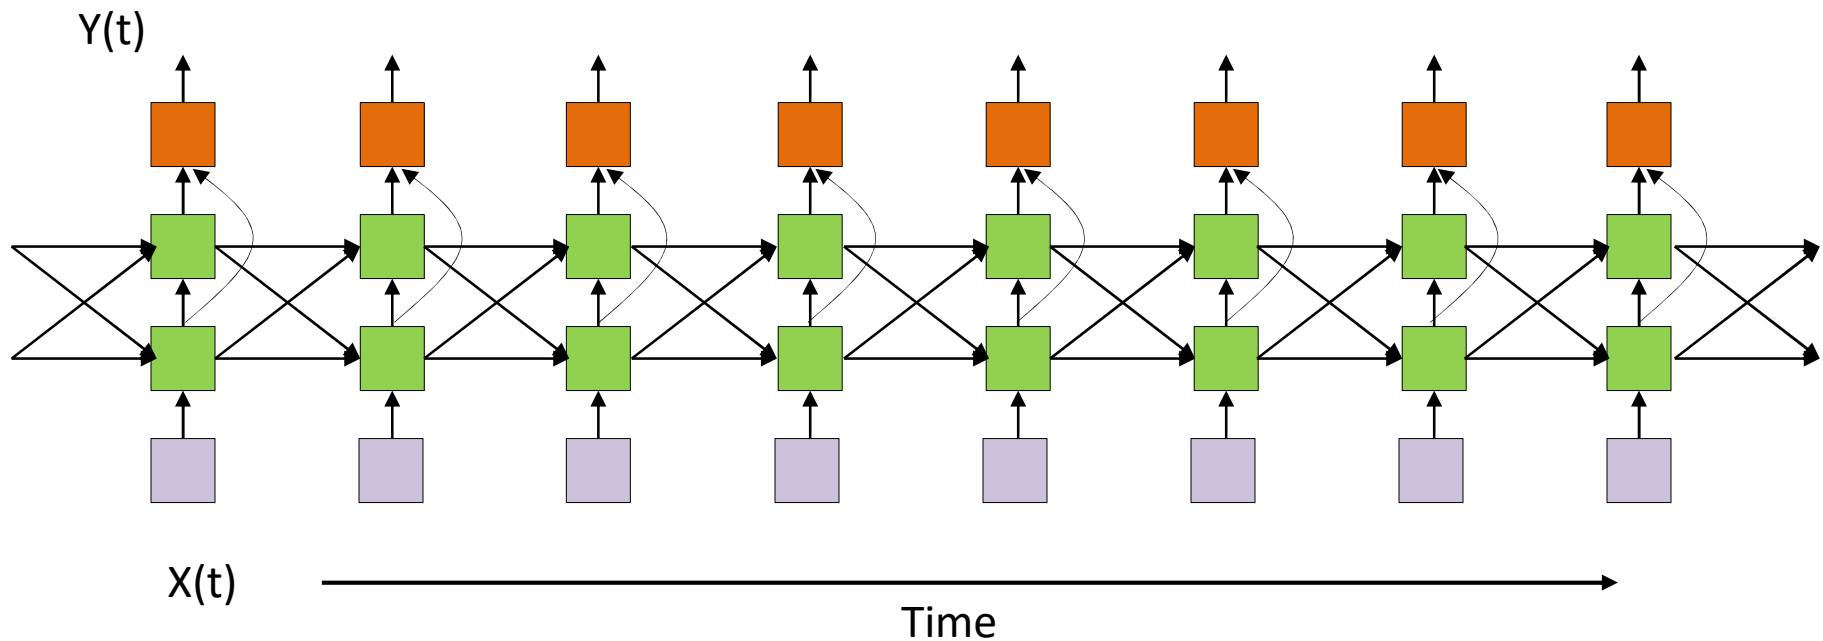
- Recurrent neural network
- All columns are identical
- *An input at $t=0$ affects outputs forever*

Multiple recurrent layer RNN



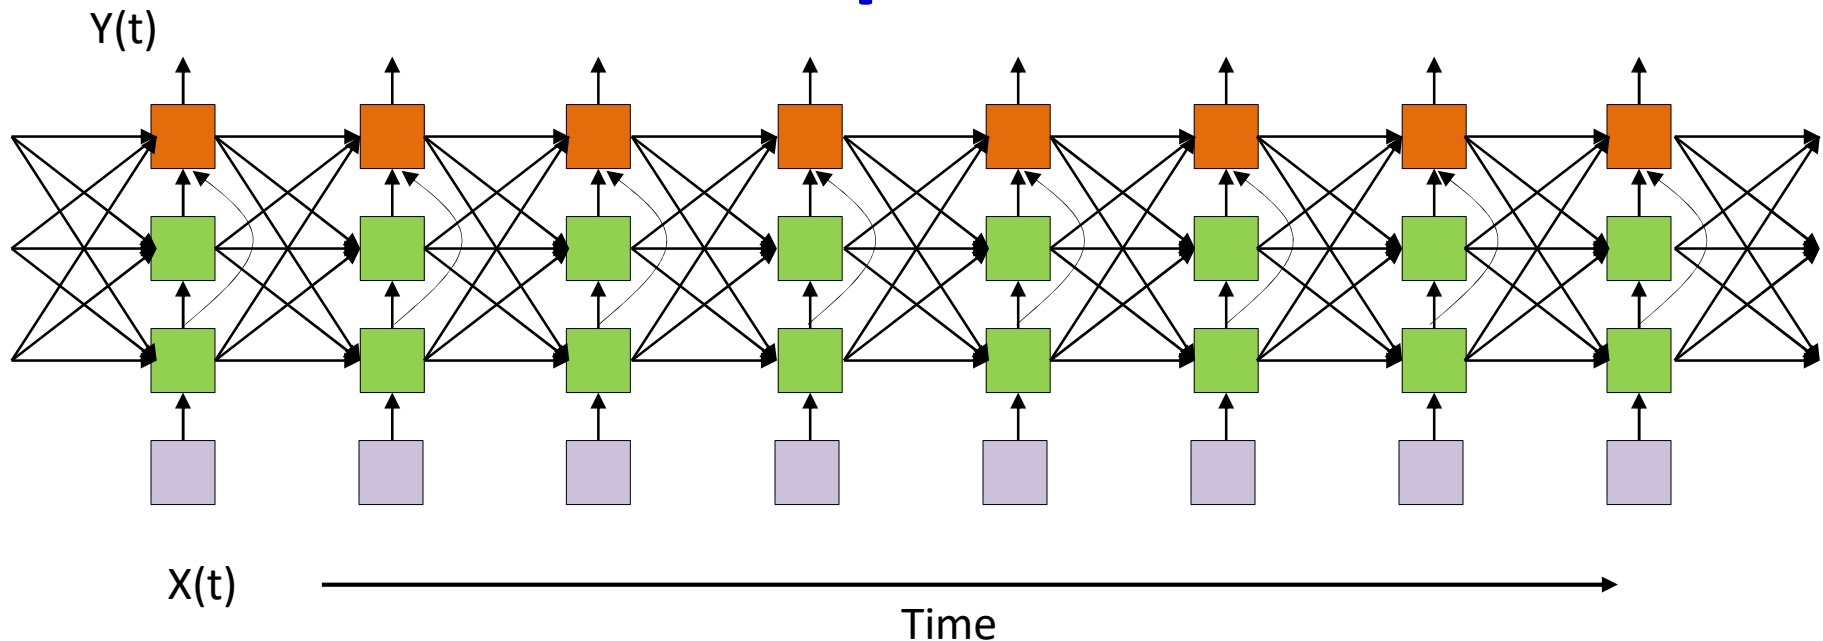
- We can also have skips..

A more complex state



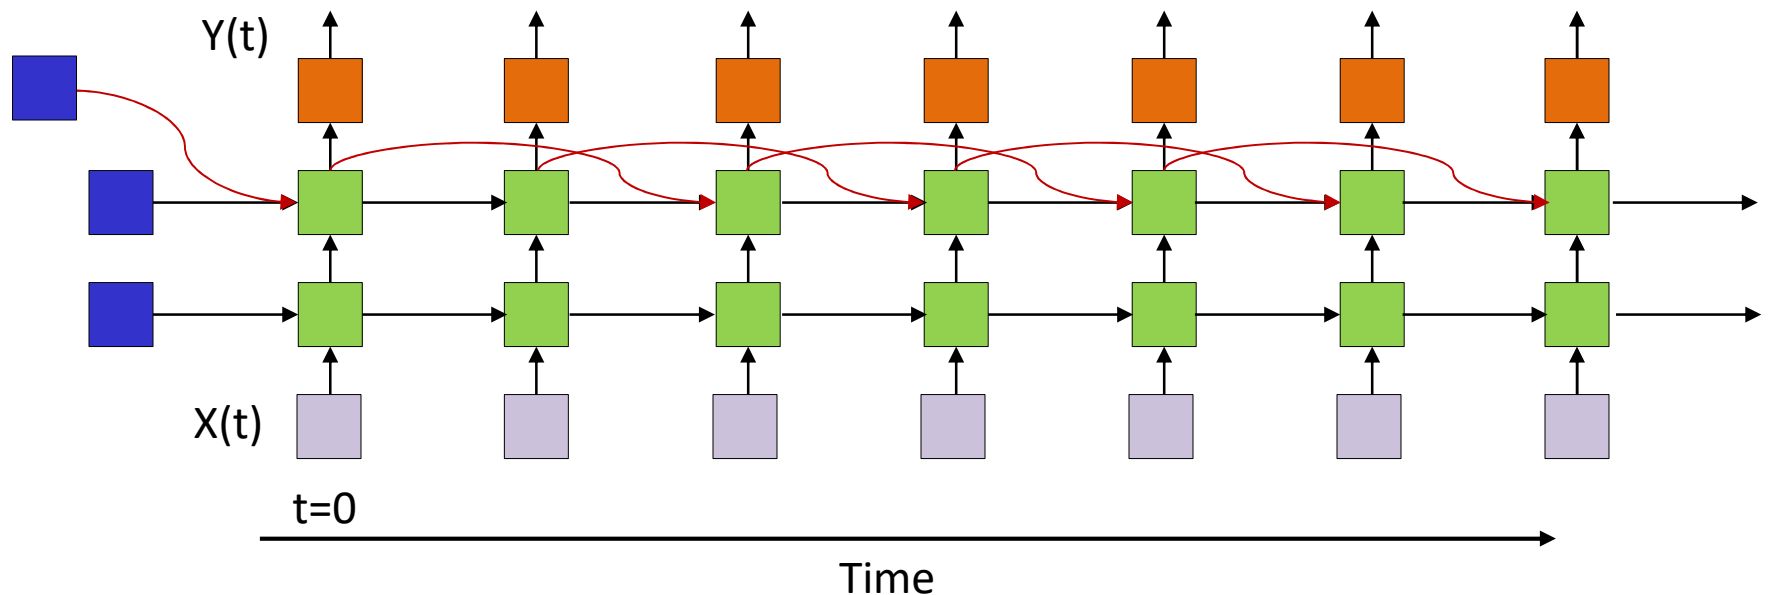
- All columns are identical
- *An input at $t=0$ affects outputs forever*

Or the network may be even more complicated




- Shades of NARX
- All columns are identical
- *An input at $t=0$ affects outputs forever*

Generalization with other recurrences




- All columns (including incoming edges) are identical

The simplest structures are most popular



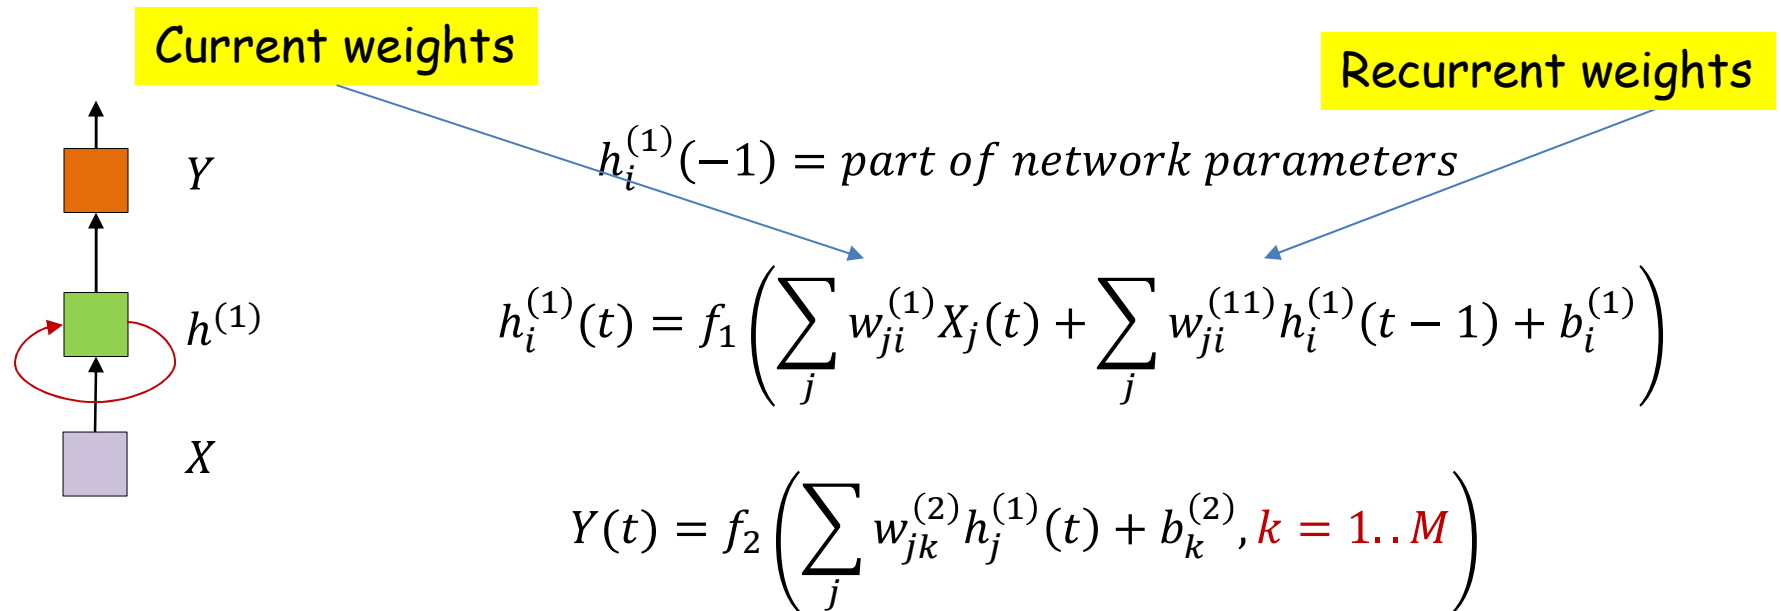
- Recurrent neural network
- All columns are identical
- *An input at $t=0$ affects outputs forever*

A Recurrent Neural Network




- Simplified models often drawn
- The loops imply recurrence

Equations



- Note superscript in indexing, which indicates layer of network from which inputs are obtained
- Assuming vector function at output, e.g. softmax
- The *state* node activation, $f_1()$ is typically $\tanh()$
- Every neuron also has a *bias* input

Equations



$h_i^{(1)}(-1) = \text{part of network parameters}$

$h_i^{(2)}(-1) = \text{part of network parameters}$

$$h_i^{(1)}(t) = f_1 \left(\sum_j w_{ji}^{(1)} X_j(t) + \sum_j w_{ji}^{(11)} h_i^{(1)}(t-1) + b_i^{(1)} \right)$$

$$h_i^{(2)}(t) = f_2 \left(\sum_j w_{ji}^{(2)} h_j^{(1)}(t) + \sum_j w_{ji}^{(22)} h_i^{(2)}(t-1) + b_i^{(2)} \right)$$


$$Y(t) = f_3 \left(\sum_j w_{jk}^{(3)} h_j^{(2)}(t) + b_k^{(3)}, k = 1..M \right)$$

- Assuming vector function at output, e.g. softmax $f_3()$
- The *state* node activations, $f_k()$ are typically $\tanh()$
- Every neuron also has a *bias* input

Equations

$h_i^{(1)}(-1) = \text{part of network parameters}$

$h_i^{(2)}(-1) = \text{part of network parameters}$




$$h_i^{(1)}(t) = f_1 \left(\sum_j w_{ji}^{(0,1)} X_j(t) + \sum_i w_{ii}^{(1,1)} h_i^{(1)}(t-1) + b_i^{(1)} \right)$$

$$h_i^{(2)}(t) = f_2 \left(\sum_j w_{ji}^{(1,2)} h_j^{(1)}(t) + \sum_j w_{ji}^{(0,2)} X_j(t) + \sum_i w_{ii}^{(2,2)} h_i^{(2)}(t-1) + b_i^{(2)} \right)$$


$$Y_i(t) = f_3 \left(\sum_j w_{jk}^{(2)} h_j^{(2)}(t) + \sum_j w_{jk}^{(1,3)} h_j^{(1)}(t) + b_k^{(3)}, k = 1..M \right)$$

Variants on recurrent nets


one to one



one to many




many to one



Images from
Karpathy

- 1: Conventional MLP
- 2: Sequence *generation*, e.g. image to caption
- 3: Sequence based *prediction or classification*, e.g. Speech recognition, text classification

Variants



- 1: *Delayed* sequence to sequence, e.g. machine translation
- 2: Sequence to sequence, e.g. stock problem, label prediction
- Etc...