

# Diffusion Models

Naresh Manwani


Machine Learning Lab, IIIT Hyderabad



INTERNATIONAL INSTITUTE OF  
INFORMATION TECHNOLOGY

H Y D E R A B A D

- 1 A diffusion model consists of an encoder and a decoder.
- 2 The encoder takes a data sample  $\mathbf{x}$  and maps it through a series of intermediate latent variables  $\mathbf{z}_1, \dots, \mathbf{z}_T$ .
- 3 The decoder reverses this process; it starts with  $\mathbf{z}_T$  and maps back through  $\mathbf{z}_{T-1}, \dots, \mathbf{z}_1$  until it recreates a data point  $\mathbf{x}$ .
- 4 In both encoder and decoder, the mappings are stochastic rather than deterministic.



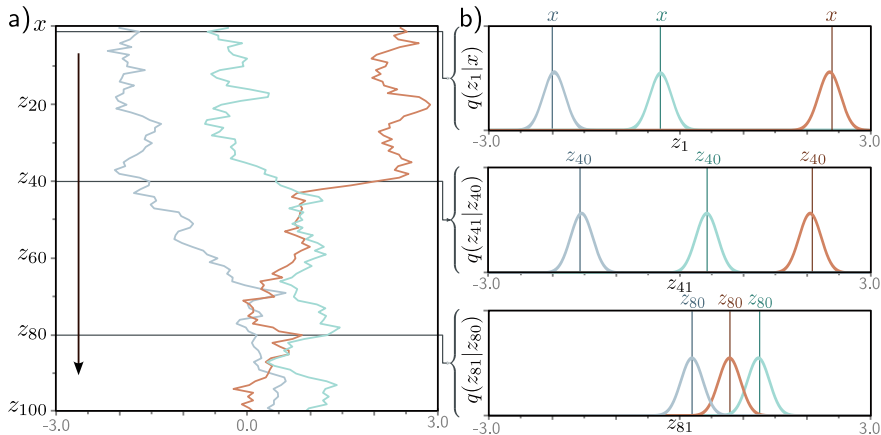
The diffusion or forward process maps a data example  $\mathbf{x}$  through a series of intermediate variables  $\mathbf{z}_1, \dots, \mathbf{z}_T$  according to:

$$\mathbf{z}_1 = \sqrt{1 - \beta_1} \mathbf{x} + \sqrt{\beta_1} \epsilon_1$$

$$\mathbf{z}_t = \sqrt{1 - \beta_t} \mathbf{z}_{t-1} + \sqrt{\beta_t} \epsilon_t, \quad t \in \{1, \dots, T\}$$

- $\epsilon_t$  is noise drawn from a standard normal distribution ( $\mathcal{N}(\mathbf{0}, I)$ ).
- The hyperparameters,  $\beta_t \in [0, 1]$ ,  $t = 1 \dots T$ , determine how quickly the noise is blended.
- $\beta_t$ ,  $t = 1 \dots T$  are collectively known as the noise schedule.

# Example: Forward Process



**Figure:**  $\beta_t = 0.03, \forall t \in \{1, \dots, 100\}$ . The conditional probabilities  $q(z_1 | x)$  and  $q(z_t | z_{t-1})$  are normal distributions with a mean that is slightly closer to zero than the current point and a fixed variance  $\beta_t$ .

The forward process can equivalently be written as:

$$q(\mathbf{z}_1|\mathbf{x}) = \mathcal{N}(\sqrt{1 - \beta_1} \mathbf{x}, \beta_1 I)$$
$$q(\mathbf{z}_t|\mathbf{z}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} \mathbf{z}_{t-1}, \beta_t I)$$

- 1 This is a Markov chain because the probability  $\mathbf{z}_t$  depends only on the value of the immediately preceding variable  $\mathbf{z}_{t-1}$ .
- 2 With sufficient steps  $T$ ,  $q(\mathbf{z}_T|\mathbf{x}) = q(\mathbf{z}_T)$  becomes a standard normal distribution.
- 3 The joint distribution of  $\mathbf{z}_1, \dots, \mathbf{z}_T$  given  $\mathbf{x}$  is as follows.

$$q(\mathbf{z}_1, \mathbf{z}_2, \dots, \mathbf{z}_T|\mathbf{x}) = q(\mathbf{z}_1|\mathbf{x}) \prod_{t=2}^T q(\mathbf{z}_t|\mathbf{z}_{t-1})$$

- ① We can see that

$$\mathbf{z}_t = \sqrt{\alpha_t} \mathbf{x} + \sqrt{1 - \alpha_t} \epsilon$$


where  $\alpha_t = \prod_{s=1}^t (1 - \beta_s)$  and  $\epsilon \sim \mathcal{N}(\mathbf{0}, I)$ .

- ② Thus,

$$q(\mathbf{z}_t|\mathbf{x}) = \mathcal{N}(\sqrt{\alpha_t} \mathbf{x}, (1 - \alpha_t)I)$$

- ③ **Advantage of Diffusion Kernel:** Having a closed form expression for  $q(\mathbf{z}_t|\mathbf{x})$  allows us to directly draw samples  $\mathbf{z}_t$  given initial data point  $\mathbf{x}$  without computing the intermediate variables  $\mathbf{z}_1, \dots, \mathbf{z}_{t-1}$ .

# Example: Diffusion Kernel



**Figure:** Started with  $x^* = 2$ . Cyan lines show  $\pm 2$  standard deviations from the mean.


# Marginal Distribution $q(\mathbf{z}_t)$ and Conditional Distribution $q(\mathbf{z}_{t-1}|\mathbf{z}_t)$



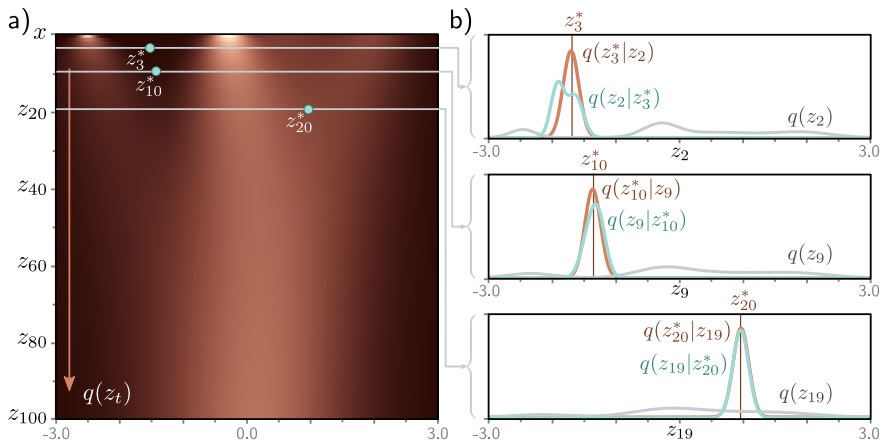
- ① **Marginal Distribution:**  $q(\mathbf{z}_t) = \int q(\mathbf{z}_t|\mathbf{x})q(\mathbf{x})d\mathbf{x}$ 
  - Cannot write closed form for  $q(\mathbf{z}_t)$  because we do not  $q(\mathbf{x})$ .
- ② **Conditional Distribution  $q(\mathbf{z}_{t-1}|\mathbf{z}_t)$ :**  $q(\mathbf{z}_{t-1}|\mathbf{z}_t) = \frac{q(\mathbf{z}_t|\mathbf{z}_{t-1})q(\mathbf{z}_{t-1})}{q(\mathbf{z}_t)}$ .
  - This is intractable since we can not compute the marginal distribution  $q(\mathbf{z}_{t-1})$ .



# Example: Marginal Diffusion Density $q(\mathbf{z}_t)$



# Example: Conditional Distribution $q(\mathbf{z}_{t-1}|\mathbf{z}_t)$





- 1 Conditional distribution  $q(\mathbf{z}_{t-1}|\mathbf{z}_t)$  is intractable since we can not compute the marginal distribution  $q(\mathbf{z}_{t-1})$ .
- 2 However, if we know the starting variable  $\mathbf{x}$ , then we do know the distribution  $q(\mathbf{z}_{t-1}|\mathbf{x})$  (this is diffusion kernel and has Gaussian density).
- 3 Thus, it is possible to find closed form of distribution  $q(\mathbf{z}_{t-1}|\mathbf{z}_t, \mathbf{x})$ .
- 4 This is used in decoder.



$$q(\mathbf{z}_{t-1}|\mathbf{z}_t, \mathbf{x})$$

$$\begin{aligned} q(\mathbf{z}_{t-1}|\mathbf{z}_t, \mathbf{x}) &= \frac{q(\mathbf{z}_t|\mathbf{z}_{t-1}, \mathbf{x})q(\mathbf{z}_{t-1}|\mathbf{x})}{q(\mathbf{z}_t|\mathbf{x})} \\ &\propto q(\mathbf{z}_t|\mathbf{z}_{t-1})q(\mathbf{z}_{t-1}|\mathbf{x}) \\ &= \mathcal{N}_{\mathbf{z}_t}(\sqrt{1-\beta_t} \mathbf{z}_{t-1}, \beta_t I) \mathcal{N}_{\mathbf{z}_{t-1}}(\sqrt{\alpha_{t-1}} \mathbf{x}, (1-\alpha_{t-1})I) \\ &\propto \mathcal{N}_{\mathbf{z}_{t-1}}\left(\frac{1}{\sqrt{1-\beta_t}} \mathbf{z}_t, \frac{\beta_t}{1-\beta_t} I\right) \mathcal{N}_{\mathbf{z}_{t-1}}(\sqrt{\alpha_{t-1}} \mathbf{x}, (1-\alpha_{t-1})I) \end{aligned}$$

where we have used the following change of variable result for Gaussian distribution.

$$\mathcal{N}_{\mathbf{x}}(A\mathbf{y} + \mathbf{b}, \Sigma) \propto \mathcal{N}_{\mathbf{y}}((A^T \Sigma A)^{-1} A^T \Sigma^{-1}(\mathbf{x} - \mathbf{b}), (A^T \Sigma A)^{-1})$$



## $q(\mathbf{z}_{t-1}|\mathbf{z}_t, \mathbf{x})$ –Continue

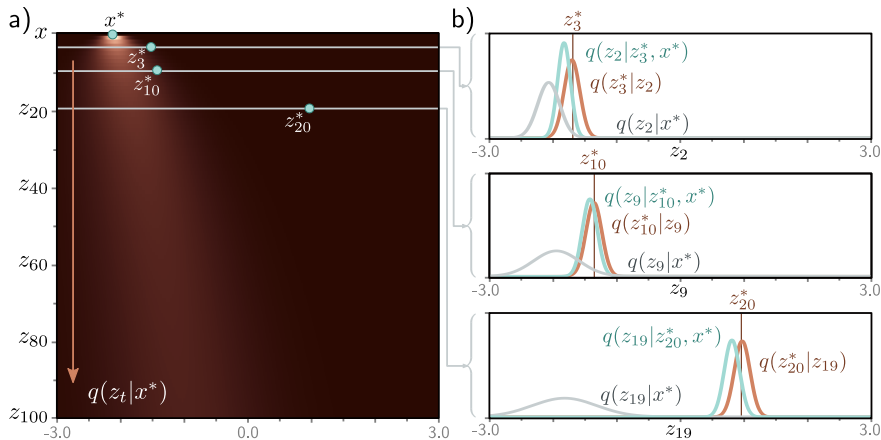
- 1 We now use the following results

$$\mathcal{N}_{\mathbf{w}}(\mathbf{a}, A)\mathcal{N}_{\mathbf{w}}(\mathbf{b}, B) \propto \mathcal{N}_{\mathbf{w}}((A^{-1} + B^{-1})^{-1})(A^{-1}\mathbf{a} + B^{-1}\mathbf{b}, (A^{-1} + B^{-1})^{-1})$$

- 2 We get the following form for  $p(\mathbf{z}_{t-1}|\mathbf{z}_t, \mathbf{x})$

$$q(\mathbf{z}_{t-1}|\mathbf{z}_t, \mathbf{x}) = \mathcal{N}_{\mathbf{z}_{t-1}} \left( \frac{1 - \alpha_{t-1}}{1 - \alpha_t} \sqrt{1 - \beta_t} \mathbf{z}_t + \frac{\sqrt{\alpha_{t-1}}}{1 - \alpha_t} \beta_t \mathbf{x}, \frac{\beta_t(1 - \alpha_{t-1})}{1 - \alpha_t} I \right)$$

# Example: Conditional Distribution $q(\mathbf{z}_{t-1}|\mathbf{z}_t, \mathbf{x})$



- 1 In decoder, we learn a series of probabilistic mappings back from latent variable  $\mathbf{z}_T$  to  $\mathbf{z}_{T-1}$ , from  $\mathbf{z}_{T-1}$  to  $\mathbf{z}_{T-2}$ , and so on, until we reach the data  $\mathbf{x}$ .
- 2 The true reverse distributions  $p(\mathbf{z}_{t-1}|\mathbf{z}_t)$  of the diffusion process are complex multi-modal distributions that depend on the data distribution  $p(\mathbf{x})$ .
- 3 We approximate these as normal distributions.

$$\begin{aligned}p(\mathbf{z}_T) &= \mathcal{N}(\mathbf{0}, I) \\p(\mathbf{z}_{t-1}|\mathbf{z}_t, \phi_t) &= \mathcal{N}(\mathbf{f}_t(\mathbf{z}_t, \phi_t), \sigma_t^2 I) \\p(\mathbf{x}|\mathbf{z}_1, \phi_1) &= \mathcal{N}(\mathbf{f}_1(\mathbf{z}_1, \phi_1), \sigma_1^2 I)\end{aligned}$$

where  $\mathbf{f}_t(\mathbf{z}_t, \phi_t)$  is a neural network that computes the mean of the normal distribution in the estimated mapping from  $\mathbf{z}_t$  to  $\mathbf{z}_{t-1}$ .

- 4 The terms  $\sigma_t^2$ ,  $t = 1 \dots T$  are predetermined.

- 1 The joint distribution of the observed variable  $\mathbf{x}$  and the latent variables  $\mathbf{z}_1, \dots, \mathbf{z}_T$  is

$$p(\mathbf{x}, \mathbf{z}_{1:T} | \phi_{1:T}) = p(\mathbf{x} | \mathbf{z}_1, \phi_1) \prod_{t=2}^T p(\mathbf{z}_{t-1} | \mathbf{z}_t, \phi_t) p(\mathbf{z}_T)$$

- 2 The likelihood of the observed data is found by marginalizing over the latent variables

$$p(\mathbf{x} | \phi_{1:T}) = \int p(\mathbf{x}, \mathbf{z}_{1:T} | \phi_{1:T}) d\mathbf{z}_{1:T} \quad (1)$$

- 3 To train the model, we maximize the log-likelihood of the training data  $\{\mathbf{x}_i\}$  with respect to  $\phi_{1:T}$

$$\hat{\phi}_{1:T} = \arg \max_{\phi_{1:T}} \sum_{i=1}^N \log p(\mathbf{x}_i | \phi_{1:T})$$

- 4 We can't maximize this directly because the marginalization in equation (1) is intractable.
- 5 We use Jensen's inequality to define a lower bound on the likelihood and optimize the parameters  $\phi_{1:T}$  with respect to this bound.





$$\begin{aligned}\log p(\mathbf{x}|\phi_{1...T}) &= \log \left[ \int p(\mathbf{x}, \mathbf{z}_{1...T}|\phi_{1...T}) d\mathbf{z}_{1...T} \right] \\ &= \log \left[ \int q(\mathbf{z}_{1...T}|\mathbf{x}) \frac{p(\mathbf{x}, \mathbf{z}_{1...T}|\phi_{1...T})}{q(\mathbf{z}_{1...T}|\mathbf{x})} d\mathbf{z}_{1...T} \right] \\ &\geq \int q(\mathbf{z}_{1...T}|\mathbf{x}) \log \left[ \frac{p(\mathbf{x}, \mathbf{z}_{1...T}|\phi_{1...T})}{q(\mathbf{z}_{1...T}|\mathbf{x})} \right] d\mathbf{z}_{1...T}\end{aligned}$$

This gives the evidence lower bound (ELBO):

$$ELBO[\phi_{1...T}] = \int q(\mathbf{z}_{1...T}|\mathbf{x}) \log \left[ \frac{p(\mathbf{x}, \mathbf{z}_{1...T}|\phi_{1...T})}{q(\mathbf{z}_{1...T}|\mathbf{x})} \right] d\mathbf{z}_{1...T}$$

- 1 In diffusion models, the decoder is trained to make the bound tighter by changing its parameters



$$\begin{aligned} ELBO[\phi_{1...T}] &= \mathbb{E}_{q(\mathbf{z}_1|\mathbf{x})} [\log p(\mathbf{x}|\mathbf{z}_1, \phi_1)] \\ &\quad - \sum_{t=2}^T \mathbb{E}_{q(\mathbf{z}_1|\mathbf{x})} [KL[q(\mathbf{z}_{t-1}|\mathbf{z}_t, \mathbf{x})||p(\mathbf{z}_{t-1}|\mathbf{z}_t, \phi_t)]] \end{aligned}$$

- The first probability term in the ELBO is

$$p(\mathbf{x}|\mathbf{z}_1, \phi_1) = \mathcal{N}_{\mathbf{z}_{t-1}}(\mathbf{f}_t[\mathbf{z}_t, \phi_t], \sigma_t^2 I)$$

The ELBO will be larger if the model prediction matches the observed data.

- $KL[q(\mathbf{z}_{t-1}|\mathbf{z}_t, \mathbf{x})||p(\mathbf{z}_{t-1}|\mathbf{z}_t, \phi_t)]$  has closed form expression. Many of the terms are independent of  $\phi_t$ .

$$\begin{aligned} & KL[q(\mathbf{z}_{t-1}|\mathbf{z}_t, \mathbf{x})||p(\mathbf{z}_{t-1}|\mathbf{z}_t, \phi_t)] \\ &= \left\| \frac{1 - \alpha_{t-1}}{1 - \alpha_t} \sqrt{1 - \beta_t} \mathbf{z}_t + \frac{\sqrt{\alpha_{t-1}\beta_t}}{1 - \alpha_t} \mathbf{x} - \mathbf{f}_t[\mathbf{z}_t, \phi_t] \right\|^2 + C \end{aligned}$$

- To fit the model, we maximize the ELBO with respect to the parameters  $\phi_{1...T}$ .
- We recast this as a minimization by multiplying with minus one and approximating the expectations with samples to give the loss function:

$$L[\phi_{1...T}] = \sum_{i=1}^I -\log(\mathcal{N}_{\mathbf{x}_i}(\mathbf{f}_1[\mathbf{z}_{i1}, \phi_1], \sigma_1^2 I)) \\ + \sum_{t=2}^T \left\| \frac{1 - \alpha_{t-1}}{1 - \alpha_t} \sqrt{1 - \beta_t} \mathbf{z}_{it} + \frac{\sqrt{\alpha_{t-1}\beta_t}}{1 - \alpha_t} \mathbf{x}_i - \mathbf{f}_t[\mathbf{z}_{it}, \phi_t] \right\|^2$$

where  $\mathbf{x}_i$  is the  $i^{th}$  data point and  $\mathbf{z}_{it}$  is the associated latent variable at  $t^{th}$  diffusion step.

- $\mathbf{f}_t[\mathbf{z}_t, \phi_t]$  predicts the value of  $\mathbf{z}_{t-1}$ .



- Diffusion models have been found to work better with a different parameterization.
- The loss function is modified so that the model aims to predict the noise that was mixed with the original data example to create the current variable.
- To achieve that, we reparameterize both the target and the network.

- The diffusion update is

$$\mathbf{z}_t = \sqrt{\alpha_t} \mathbf{x} + \sqrt{1 - \alpha_t} \epsilon$$

- It follows that the data term  $\mathbf{x}$  can be expressed as the diffused image  $\mathbf{z}_t$  minus the noise that was added to it.

$$\mathbf{x} = \frac{1}{\sqrt{\alpha_t}} \mathbf{z}_t - \frac{\sqrt{1 - \alpha_t}}{\sqrt{\alpha_t}} \epsilon$$

- Using this, the target is modified as

$$\frac{1 - \alpha_{t-1}}{1 - \alpha_t} \sqrt{1 - \beta_t} \mathbf{z}_t + \frac{\sqrt{\alpha_{t-1}} \beta_t}{1 - \alpha_t} \mathbf{x} = \frac{1}{\sqrt{1 - \beta_t}} \mathbf{z}_t - \frac{\beta_t}{\sqrt{1 - \alpha_t} \sqrt{1 - \beta_t}} \epsilon$$

- We replace the model  $\hat{\mathbf{z}}_{t-1} = \mathbf{f}_t[\mathbf{z}_t, \phi_t]$  as

$$\mathbf{f}_t[\mathbf{z}_t, \phi_t] = \frac{1}{\sqrt{1 - \beta_t}} \mathbf{z}_t - \frac{\beta_t}{\sqrt{1 - \alpha_t} \sqrt{1 - \beta_t}} \mathbf{g}_t[\mathbf{z}_t, \phi_t]$$

where  $\mathbf{g}_t[\mathbf{z}_t, \phi_t]$  tries to approximate the noise  $\epsilon$  that was added to  $\mathbf{x}$  to create  $\mathbf{z}_t$ .



$$\begin{aligned} L[\phi_{1...T}] &= \sum_{i=1}^I \sum_{t=1}^T \|\mathbf{g}_t[\mathbf{z}_{it}, \phi_t] - \epsilon_{it}\|^2 \\ &= \sum_{i=1}^I \sum_{t=1}^T \|\mathbf{g}_t[\sqrt{\alpha_t} \mathbf{x}_i + \sqrt{1 - \alpha_t} \epsilon_{it}, \phi_t] - \epsilon_{it}\|^2 \end{aligned}$$



---

**Algorithm 18.1:** Diffusion model training

---

**Input:** Training data  $\mathbf{x}$

**Output:** Model parameters  $\phi_t$

**repeat**

```
    for  $i \in \mathcal{B}$  do                                // For every training example index in batch
         $t \sim \text{Uniform}[1, \dots, T]$                 // Sample random timestep
         $\epsilon \sim \text{Norm}[\mathbf{0}, \mathbf{I}]$                 // Sample noise
         $\ell_i = \left\| \mathbf{g}_t \left[ \sqrt{\alpha_t} \mathbf{x}_i + \sqrt{1 - \alpha_t} \epsilon, \phi_t \right] - \epsilon \right\|^2$  // Compute individual loss
```

Accumulate losses for batch and take gradient step

**until** converged

---

- The training algorithm is simple to implement.
- It naturally augments the dataset. We can reuse every original data point  $\mathbf{x}_i$  as many times as we want at each time step with different noise instantiations  $\epsilon$ .



---

**Algorithm 18.2:** Sampling

---

**Input:** Model,  $\mathbf{g}_t[\bullet, \phi_t]$


**Output:** Sample,  $\mathbf{x}$

```
 $\mathbf{z}_T \sim \text{Norm}_{\mathbf{z}}[\mathbf{0}, \mathbf{I}]$  // Sample last latent variable
for  $t = T \dots 2$  do
     $\hat{\mathbf{z}}_{t-1} = \frac{1}{\sqrt{1-\beta_t}} \mathbf{z}_t - \frac{\beta_t}{\sqrt{1-\alpha_t}\sqrt{1-\beta_t}} \mathbf{g}_t[\mathbf{z}_t, \phi_t]$  // Predict previous latent variable
     $\epsilon \sim \text{Norm}_{\epsilon}[\mathbf{0}, \mathbf{I}]$  // Draw new noise vector
     $\mathbf{z}_{t-1} = \hat{\mathbf{z}}_{t-1} + \sigma_t \epsilon$  // Add noise to previous latent variable
 $\mathbf{x} = \frac{1}{\sqrt{1-\beta_1}} \mathbf{z}_1 - \frac{\beta_1}{\sqrt{1-\alpha_1}\sqrt{1-\beta_1}} \mathbf{g}_1[\mathbf{z}_1, \phi_1]$  // Generate sample from  $\mathbf{z}_1$  without noise
```


---

- The sampling algorithm requires serial processing of many neural networks  $\mathbf{g}_t[\mathbf{z}_t, \phi_t]$  and is hence time-consuming.

- Here, we need to construct models that can take a noisy image and predict the noise that was added at each step.
- The architectural choice for this image-to-image mapping is the U-Net.



- There may be a very large number of diffusion steps, and training and storing multiple U-Nets is inefficient.
- The solution is to train a single U-Net that also takes a predetermined vector representing the time step as input.



# Training Diffusion UNet

- A large number of time steps are needed as the conditional probabilities  $q(\mathbf{z}_{t-1}|\mathbf{z}_t)$  become closer to normal when the hyperparameters  $\beta_t$  are close to zero, matching the form of the decoder distributions  $p(\mathbf{z}_{t-1}|\mathbf{z}_t, \phi_t)$ .
- However, this makes sampling slow.
- We might have to run the U-Net model through  $T = 1000$  steps to generate good images.